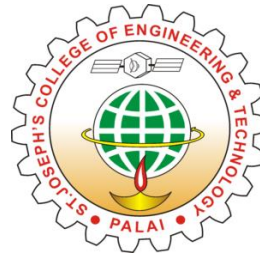


ST. JOSEPH'S  
COLLEGE OF ENGINEERING  
AND TECHNOLOGY,  
- PALAI -



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

CST 204 - Database Management Systems

**Prof. Sarju S**

30 September 2021

# CST 204 – Database Management Systems

# Module 5

---



- ▶ Transaction Processing Concepts - overview of concurrency control, Transaction Model, Significance of concurrency Control & Recovery, Transaction States, System Log, Desirable Properties of transactions.
- ▶ Serial schedules, Concurrent and Serializable Schedules, Conflict equivalence and conflict serializability, Recoverable and cascade-less schedules, Locking, Two-phase locking and its variations. Log-based recovery, Deferred database modification, check-pointing.
- ▶ Introduction to NoSQL Databases, Main characteristics of Key-value DB (examples from: Redis), Document DB (examples from: MongoDB)
- ▶ Main characteristics of Column - Family DB (examples from: Cassandra) and Graph DB (examples from : ArangoDB)

# Schedules

# Schedules

---



- ▶ A schedule (or history)  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is an ordering of the operations of the transactions.
- ▶ Operations from different transactions can be interleaved in the schedule  $S$ .

# Recoverable Schedules

- ▶ Schedules in which transactions commit only after all transactions whose changes they read commit are called recoverable schedules.
- ▶ If some transaction  $T_j$  is reading value updated or written by some other transaction  $T_i$ , then the commit of  $T_j$  must occur after the commit of  $T_i$ .

$T_1$	$T_2$
R(A)	
W(A)	
	W(A)
	R(A)
commit	
	commit

This is a recoverable schedule since  $T_1$  commits before  $T_2$ , that makes the value read by  $T_2$  correct.

# Irrecoverable Schedule

T1	T1's buffer space	T2	T2's Buffer Space	Database
				A=5000
R(A);	A=5000			A=5000
A=A-100;	A=4000			A=5000
W(A);	A=4000			A=4000
		R(A);	A=4000	A=4000
		A=A+500;	A=4500	A=4000
		W(A);	A=4500	A=4500
		Commit;		
Failure Point				
Commit;				

- ▶ The table shows a schedule with two transactions,  $T_1$  reads and writes A and that value is read and written by  $T_2$ .
- ▶  $T_2$  commits. But later on,  $T_1$  fails.
- ▶ So we have to rollback  $T_1$ . Since  $T_2$  has read the value written by  $T_1$ , it should also be rolled back.
- ▶ But we have already committed that.
- ▶ So this schedule is **irrecoverable schedule**.
- ▶ When  $T_j$  is reading the value updated by  $T_i$  and  $T_j$  is committed before committing of  $T_i$ , the schedule will be irrecoverable.

# Recoverable with Cascading Rollback:

- ▶ Cascading Rollback (or cascading abort) to occur in some recoverable schedules, where an *uncommitted transaction has to be rolled back because it read an item from a transaction that failed*.

T1	T1's buffer space	T2	T2's Buffer Space	Database
				A=5000
R(A);	A=5000			A=5000
A=A-100;	A=4000			A=5000
W(A);	A=4000			A=4000
		R(A);	A=4000	A=4000
		A=A+500;	A=4500	A=4000
		W(A);	A=4500	A=4500
Failure Point				
Commit;				
		Commit;		

- ▶ The table shows a schedule with two transactions,  $T_1$  reads and writes A and that value is read and written by  $T_2$ .
- ▶ But later on,  $T_1$  fails. So we have to rollback  $T_1$ .
- ▶ Since  $T_2$  has read the value written by  $T_1$ , it should also be rolled back.
- ▶ As it has not committed, we can rollback  $T_2$  as well. So it is **recoverable with cascading rollback**.
- ▶ *If  $T_j$  is reading value updated by  $T_i$  and commit of  $T_j$  is delayed till commit of  $T_i$ , the schedule is called recoverable with cascading rollback.*



# Cascadeless Recoverable Rollback

- ▶ A schedule is said to be cascadeless, or to avoid cascading rollback, *if every transaction in the schedule reads only items that were written by committed transactions.*

T1	T1's buffer space	T2	T2's Buffer Space	Database
				A=5000
R(A);	A=5000			A=5000
A=A-100;	A=4000			A=5000
W(A);	A=4000			A=4000
Commit;				
		R(A);	A=4000	A=4000
		A=A+500;	A=4500	A=4000
		W(A);	A=4500	A=4500
		Commit;		

- ▶ The table shows a schedule with two transactions, T1 reads and writes A and commits and that value is read by T2.
- ▶ But if T1 fails before commit, no other transaction has read its value, so there is no need to rollback other transaction.
- ▶ So this is a **Cascadeless** recoverable schedule.
- ▶ *If Tj reads value updated by Ti only after Ti is committed, the schedule will be cascadeless recoverable.*

# Serial Schedules

- ▶ Schedules in which the transactions are executed non-interleaved
  - ▶ a serial schedule is one in which no transaction starts until a running transaction has ended are called serial schedules.

$T_1$	$T_2$
R(A)	
W(A)	
R(B)	
	W(B)
	R(A)
	R(B)

Example: Consider the schedule involving two transactions  $T_1$  and  $T_2$ . This is a serial schedule since the transactions perform serially in the order  $T_1 \rightarrow T_2$



# Non-Serial Schedule

---

- ▶ This is a type of Scheduling where the operations of multiple transactions are interleaved.
  - ▶ Unlike the serial schedule where one transaction must wait for another to complete all its operation, in the non-serial schedule, the other transaction proceeds without waiting for the previous transaction to complete.
- ▶ This might lead to a rise in the concurrency problem.
- ▶ It can be of two types namely, Serializable and Non-Serializable Schedule.



# Serializable Schedule

---

- ▶ This is used to maintain the consistency of the database.
- ▶ It is mainly used in the Non-Serial scheduling to verify whether the scheduling will lead to any inconsistency or not.
- ▶ These are of two types:
  - ▶ Conflict Serializable
  - ▶ View Serializable:



# Conflict Serializable Schedule

---

- ▶ A schedule is called conflict serializability if after swapping of non-conflicting operations, it can transform into a serial schedule.
- ▶ The schedule will be a conflict serializable if it is conflict equivalent to a serial schedule.



# Conflict Serializable Schedule


---

## Conflicting Operations

- ▶ The two operations become conflicting if all conditions satisfy:
  - ▶ Both belong to separate transactions.
  - ▶ They have the same data item.
  - ▶ They contain at least one write operation.

# Conflict Serializable Schedule


## Conflicting Operations



Here,  $S1 = S2$ . That means it is non-conflict.

**Schedule S1**

**Schedule S2**



Here,  $S1 \neq S2$ . That means it is conflict

**Schedule S1**

**Schedule S2**



# Conflict Serializable Schedule

---

## Conflict Equivalent

- ▶ In the conflict equivalent, one can be transformed to another by swapping non-conflicting operations.
- ▶ Two schedules are said to be conflict equivalent if and only if:
  - ▶ They contain the same set of the transaction.
  - ▶ If each pair of conflict operations are ordered in the same way.




# Conflict Serializable Schedule

## Conflict Equivalent

T1	T2
Read(A) Write(A)	
	Read(A) Write(A)
Read(B) Write(B)	
	Read(B) Write(B)

**Schedule S1**

T1	T2
Read(A) Write(A) Read(B) Write(B)	
	Read(A) Write(A) Read(B) Write(B)



Schedule S1 can be transformed into a serial schedule by swapping non-conflicting operations of S1.

# Testing for Conflict Serializability of a Schedule




- ▶ There is a simple algorithm for determining whether a particular schedule is conflict serializable or not.
- ▶ The Algorithm can be written as:
  1. Create a node  $T$  in the graph for each participating transaction in the schedule.
  2. For the conflicting operation `read_item(X)` and `write_item(X)` – If a Transaction  $T_j$  executes a `read_item(X)` after  $T_i$  executes a `write_item(X)`, draw an edge from  $T_i$  to  $T_j$  in the graph.
  3. For the conflicting operation `write_item(X)` and `read_item(X)` – If a Transaction  $T_j$  executes a `write_item(X)` after  $T_i$  executes a `read_item(X)`, draw an edge from  $T_i$  to  $T_j$  in the graph.
  4. For the conflicting operation `write_item(X)` and `write_item(X)` – If a Transaction  $T_j$  executes a `write_item(X)` after  $T_i$  executes a `write_item(X)`, draw an edge from  $T_i$  to  $T_j$  in the graph.
  5. The Schedule  $S$  is serializable if there is no cycle in the precedence graph.

# Testing for Conflict Serializability of a Schedule - Example



T1	T2	T3
R(x)		R(y) R(x)
	R(y) R(z)	W(y)
R(z) W(x) W(z)	W(z)	

- ▶ Draw Edge when we have
- ▶ read\_item(X) and write\_item(X)
- ▶ write\_item(X) and read\_item(X)
- ▶ write\_item(X) and write\_item(X)



- ▶ T1-R(x): no W(x) in T2, T3
- ▶ T3-R(y): no W(y) in T2, T1
- ▶ T3-R(x): W(x) in T1 draw the edge T3->T1
- ▶ T2-R(y): W(y) in T3 draw the edge T2->T3
- ▶ T2-R(z): W(z) in T1 draw the edge T2->T1
- ▶ T3-W(y): no R(y) or W(y) in T2, T1
- ▶ T2-W(z): R(z) and W(z) in T1 draw edge T2->T1(already there)


- ▶ As we have no Cycle/Loop in the Precedence Graph these schedule is conflict serializable

# Testing for Conflict Serializability of a Schedule – Example 2



T1	T2	T3
R(A)	W(A)	W(A)
W(A)		

- ▶ Draw Edge when we have
- ▶ read\_item(X) and write\_item(X)
- ▶ write\_item(X) and read\_item(X)
- ▶ write\_item(X) and write\_item(X)




- ▶ As we have Cycle/Loop in the Precedence Graph this schedule is Non conflict serializable . Is this serializable?

# University Previous Question Paper Question



Check if the following schedules are conflict-serializable using precedence graph. If so, give the equivalent serial schedule(s).  $r_3(X)$ ,  $r_2(X)$ ,  $w_3(X)$ ,  $r_1(X)$ ,  $w_1(X)$ . (Note:  $r_i(X)/w_i(X)$  means transaction  $T_i$  issues read/write on item  $X$ .)

T1	T2	T3
$R(x)$ $W(x)$	$R(x)$	$R(x)$ $W(x)$



- ▶ Draw Edge when we have
- ▶  $read\_item(X)$  and  $write\_item(X)$
- ▶  $write\_item(X)$  and  $read\_item(X)$
- ▶  $write\_item(X)$  and  $write\_item(X)$

- ▶ As we have no Cycle/Loop in the Precedence Graph this schedule is conflict serializable .

# What is View Serializability?

- ▶ View Serializability is a process to find out that a given schedule is view serializable or not.
- ▶ To check whether a given schedule is view serializable, we need to check whether the given schedule is View Equivalent to its serial schedule.
- ▶ Lets take an example to understand what I mean by that.

T1	T2	T3
R(A)	W(A)	
W(A)		
		W(A)

View Equivalent

T1	T2	T3
R(A) W(A)	W(A)	
		W(A)

# View Equivalent

- ▶ Two schedules S1 and S2 are said to be view equivalent if they satisfy the following conditions:
  - ▶ Initial Read
  - ▶ An initial read of both schedules must be the same.
  - ▶ Suppose two schedule S1 and S2. In schedule S1, if a transaction T1 is reading the data item A, then in S2, transaction T1 should also read A.

T1	T2
Read(A)	Write(A)

**Schedule S1**

T1	T2
Read(A)	Write(A)

**Schedule S2**

These schedules are view equivalent because Initial read operation in S1 is done by T1 and in S2 it is also done by T1.

# View Serializability - View Equivalent

- ▶ Two schedules S1 and S2 are said to be view equivalent if they satisfy the following conditions:
  - ▶ Updated Read
    - ▶ In schedule S1, if  $T_i$  is reading A which is updated by  $T_j$  then in S2 also,  $T_i$  should read A which is updated by  $T_j$ .

T1	T2	T3
Write(A)	Write(A)	Read(A)

**Schedule S1**

T1	T2	T3
Write(A)	Write(A)	<u>Read(A)</u>

**Schedule S2**

These two schedules are not view equal because, in S1, T3 is reading A updated by T2 and in S2, T3 is reading A updated by T1.



# View Serializability - View Equivalent

- ▶ Two schedules S1 and S2 are said to be view equivalent if they satisfy the following conditions:
  - ▶ Final Write
    - ▶ A final write must be the same between both the schedules. In schedule S1, if a transaction T1 updates A at last then in S2, final writes operations should also be done by T1.

T1	T2	T3
Write(A)	Read(A)	
		Write(A)

**Schedule S1**

T1	T2	T3
Write(A)	Read(A)	
		Write(A)

**Schedule S2**

These two schedules is view equal because Final write operation in S1 is done by T3 and in S2, the final write operation is also done by T3.

# View Serializability - Example

Non-Serial

-----  
S1

-----  
T1      T2

-----  
R(X)  
W(X)

R(X)  
W(X)

R(Y)  
W(Y)

R(Y)  
W(Y)

Serial

-----  
S2

-----  
T1      T2

-----  
R(X)  
W(X)  
R(Y)  
W(Y)

R(X)  
W(X)  
R(Y)  
W(Y)

*Beginnerbook.com*

S2 is the serial schedule of S1. If we can prove that they are view equivalent then we can say that given schedule S1 is view Serializable



# View Serializability - View Equivalent

---

- ▶ Lets check the three conditions of view serializability:
- ▶ Initial Read
  - ▶ In schedule S1, transaction T1 first reads the data item X.
  - ▶ In S2 also transaction T1 first reads the data item X.
  - ▶ Lets check for Y. In schedule S1, transaction T1 first reads the data item Y. In S2 also the first read operation on Y is performed by T1.
  - ▶ We checked for both data items X & Y and the **initial read condition is satisfied** in S1 & S2.



# View Serializability - View Equivalent

---

- ▶ Lets check the three conditions of view serializability:
- ▶ Final Write
  - ▶ In schedule S1, the final write operation on X is done by transaction T2. In S2 also transaction T2 performs the final write on X.
  - ▶ Lets check for Y. In schedule S1, the final write operation on Y is done by transaction T2. In schedule S2, final write on Y is done by T2.
  - ▶ We checked for both data items X & Y and the **final write condition is satisfied** in S1 & S2.



# View Serializability - View Equivalent

---

- ▶ Lets check the three conditions of view serializability:
- ▶ Update Read
  - ▶ In S1, transaction T2 reads the value of X, written by T1. In S2, the same transaction T2 reads the X after it is written by T1.
  - ▶ In S1, transaction T2 reads the value of Y, written by T1. In S2, the same transaction T2 reads the value of Y after it is updated by T1.
  - ▶ The **update read condition is also satisfied** for both the schedules.
- ▶ **Result:** Since all the three conditions that checks whether the two schedules are view equivalent are satisfied in this example, which means S1 and S2 are view equivalent.

# References

---



- ▶ <https://www.geeksforgeeks.org/recoverability-in-dbms/>
- ▶ <https://www.geeksforgeeks.org/types-of-schedules-in-dbms/>
- ▶ <https://www.javatpoint.com/dbms-conflict-serializable-schedule>
- ▶ <https://beginnersbook.com/2018/12/dbms-view-serializability/>



# Thank You



**Prof. Sarju S**

Department of Computer Science and Engineering  
St. Joseph's College of Engineering and Technology, Palai  
sarju.s@sjcetpalai.ac.in  sarju-s