

## Concurrency Control Protocols

- Aims: Achieve Sensitivity -  
- " Recoverability -

Earlier we studied to check if schedule is sensigable or not.  
But in CED, we study how to make our schedule sensigable. (It's not to check)  
↳ How to make a sensible, and recoverable

How to achieve this  $\Rightarrow$  By using locking protocols

① Shared exclusive locking

② Shared lock. - Read only

③ Exclusive lock - Read / write.

Eg  $T_1$   $T_2$ .

$S(A) \leftarrow \text{lock}$	$X(A) \leftarrow \text{lock}$
$R(A)$	$R(A)$
$U(A) \leftarrow \text{unlock}$	$W(A)$
$U(A) \leftarrow \text{unlock}$	

only read . Both read & write .

① granted ( $S(A)$ )

Now Request  $S(A) \rightarrow$  OK

Request

S	X
✓	X
X	X

② granted  $S(A)$

Request -  $X(A)$  - No of Read Write Conflict?

Shared - Ex lock does not guarantee  
serializable always.

① May not sufficient to produce serializable.

$\rightarrow T_1 \rightarrow T_2$

$T_1 \rightarrow T_2$  or  
 $T_2 \rightarrow T_1$ .

T1	T2
R(A) W(A)	
R(B) W(B)	

T1	T2
X(A) R(A) W(A) U(A)	
S(A)	R(B) U(B)

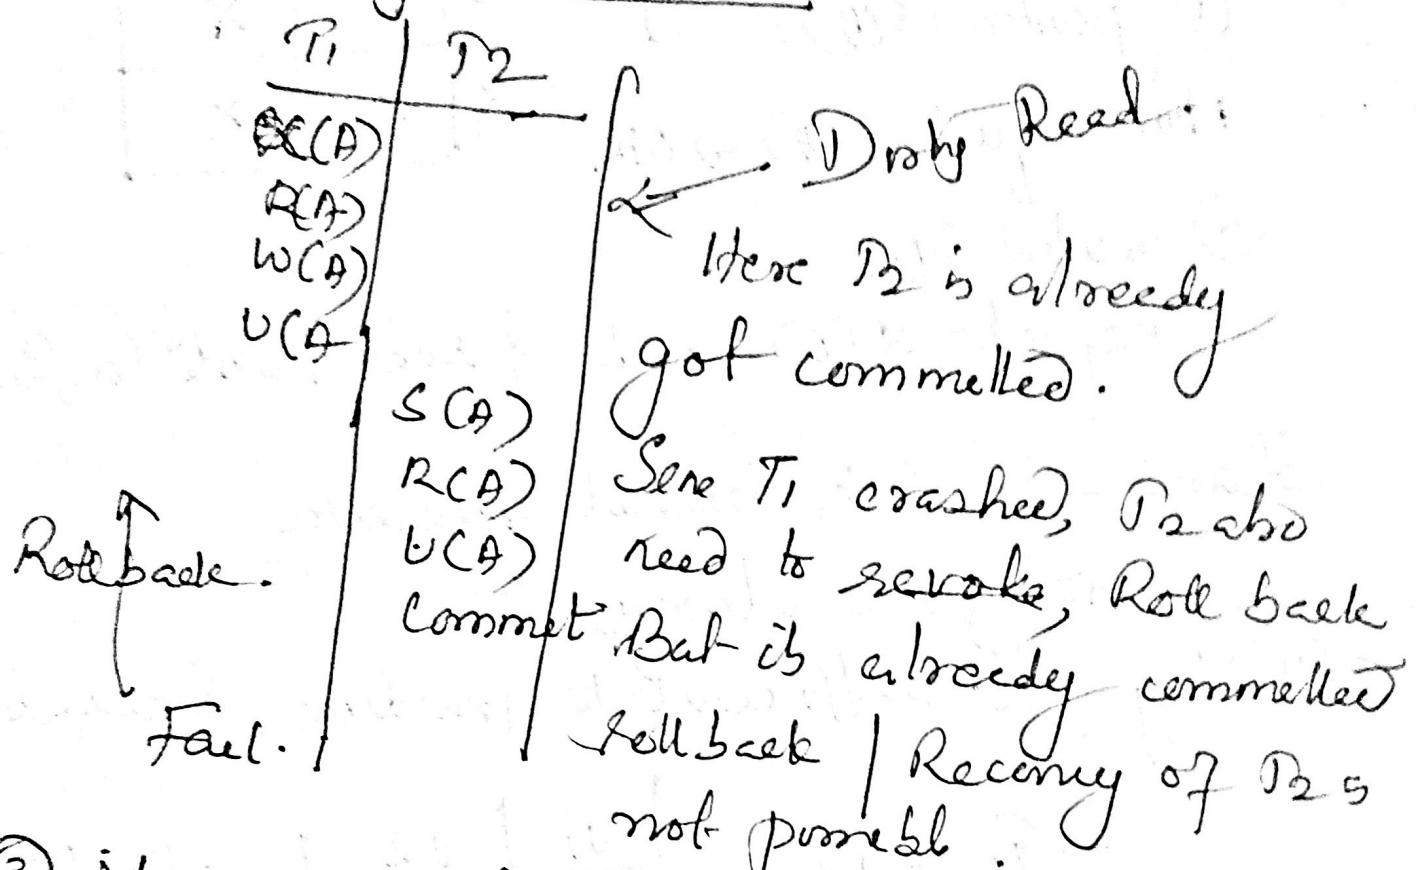
$T_1 \rightarrow T_2$

←

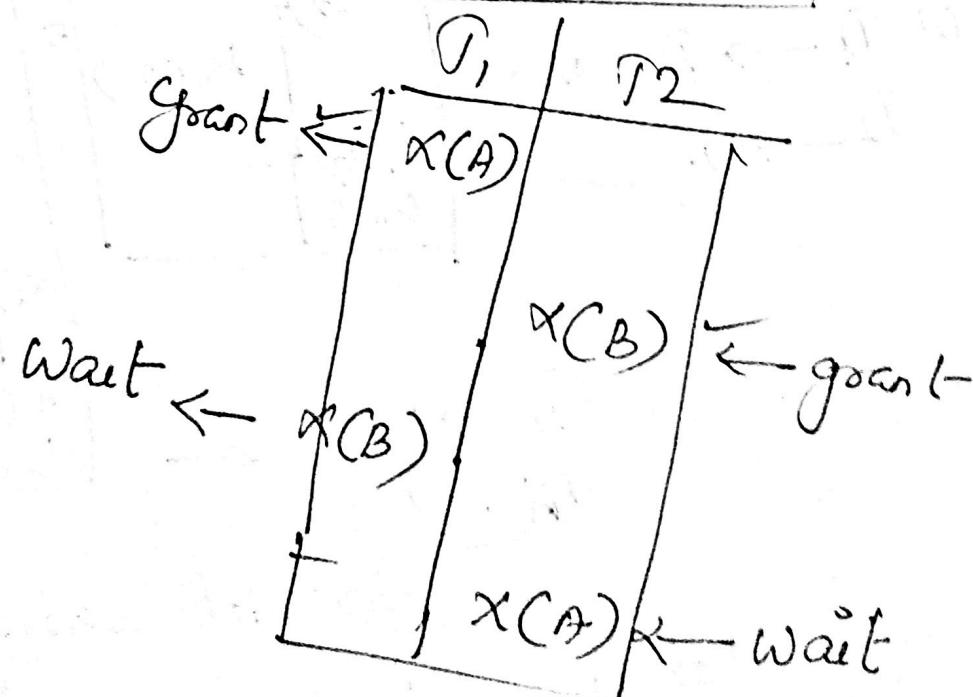
→ Not serializable.

⇒ May not sufficient to serializable

## ② Not always Recoverable.

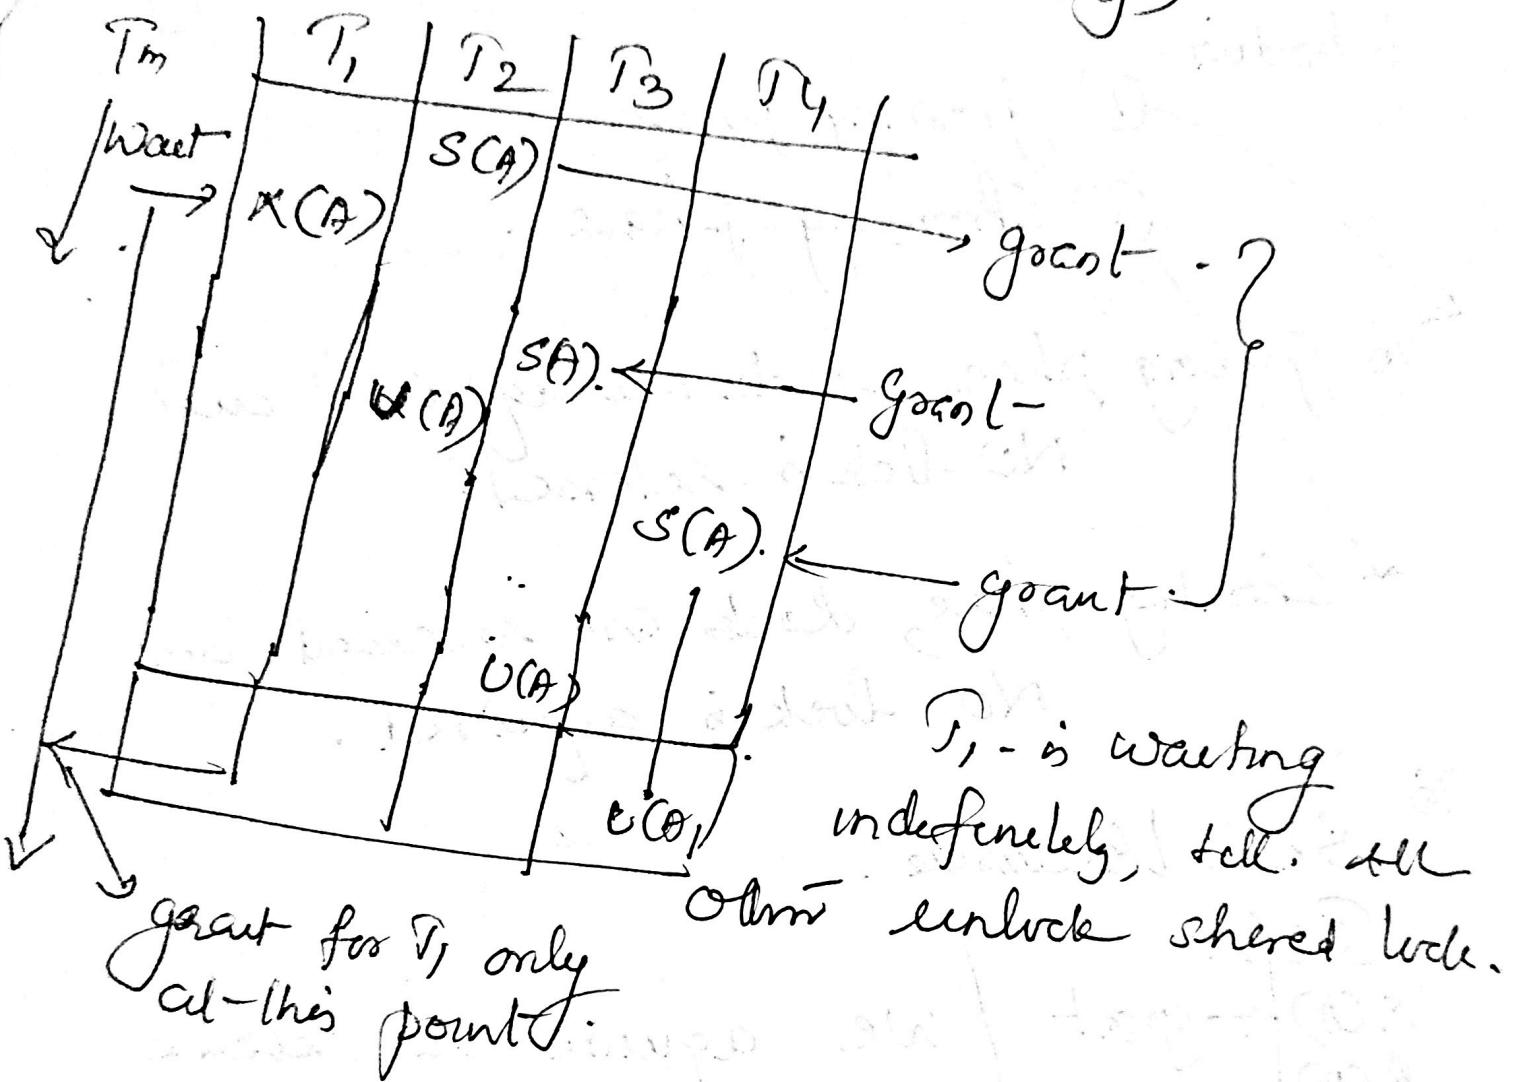


## ③ May not free from Deadlock.



(Deadlock  $\Rightarrow$  infinite waiting.)

⑤ May not free from starvation  
 (Not infinite waiting)



## 2 Phase Locking Protocol 2 PL

Show  
all

Extensions of Simple Shared / Exclusive locking discipline

- ① Growing phase
- ② Shrinking phase

1) Growing phase, locks are acquired and NO lock is released.

2) Shrinking phase, locks are released and NO lock is acquired.

### Shared / Exclusive

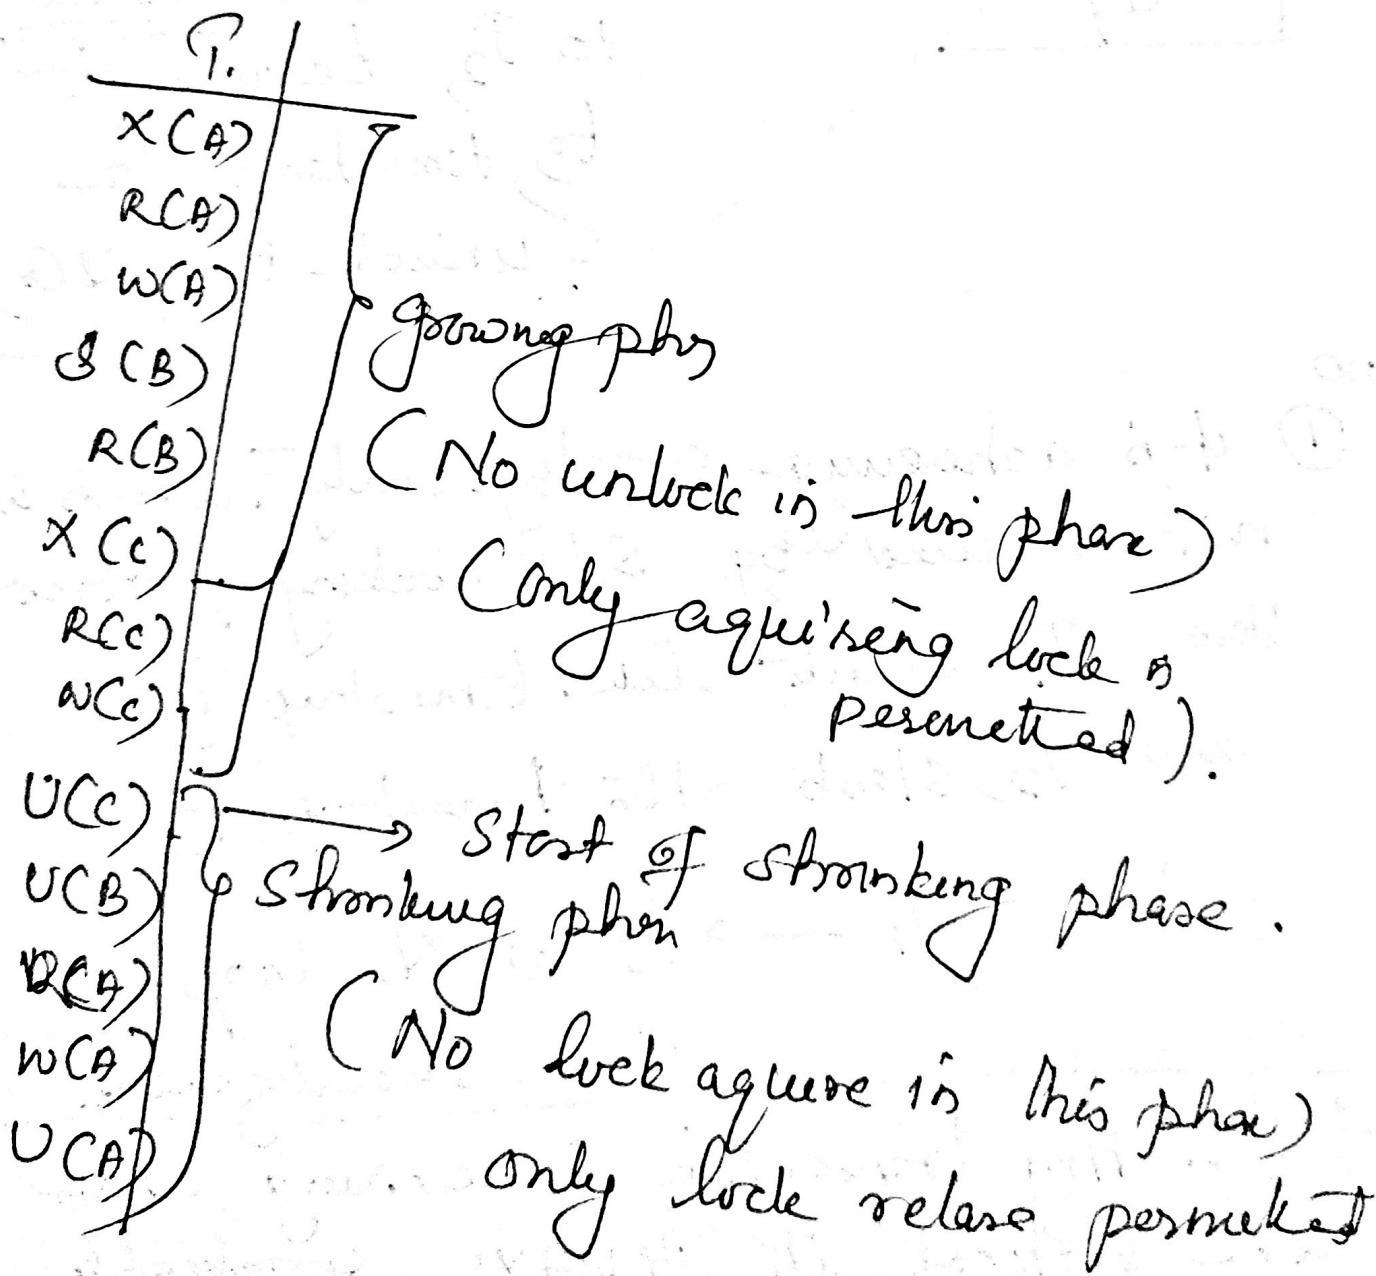
P, f.  
 $\alpha(A)$  ← grant  
 $\alpha(C)$   
 $w(A)$   
 $s(B)$  ← grant  
 $r(B)$   
 $\alpha(C)$  ← grant  
 $r(C)$   
 $w(\alpha)$   
 $v(\alpha)$  ← Rel.  
 $w(C)$   
 $v(C)$  ends when the first release occurs.

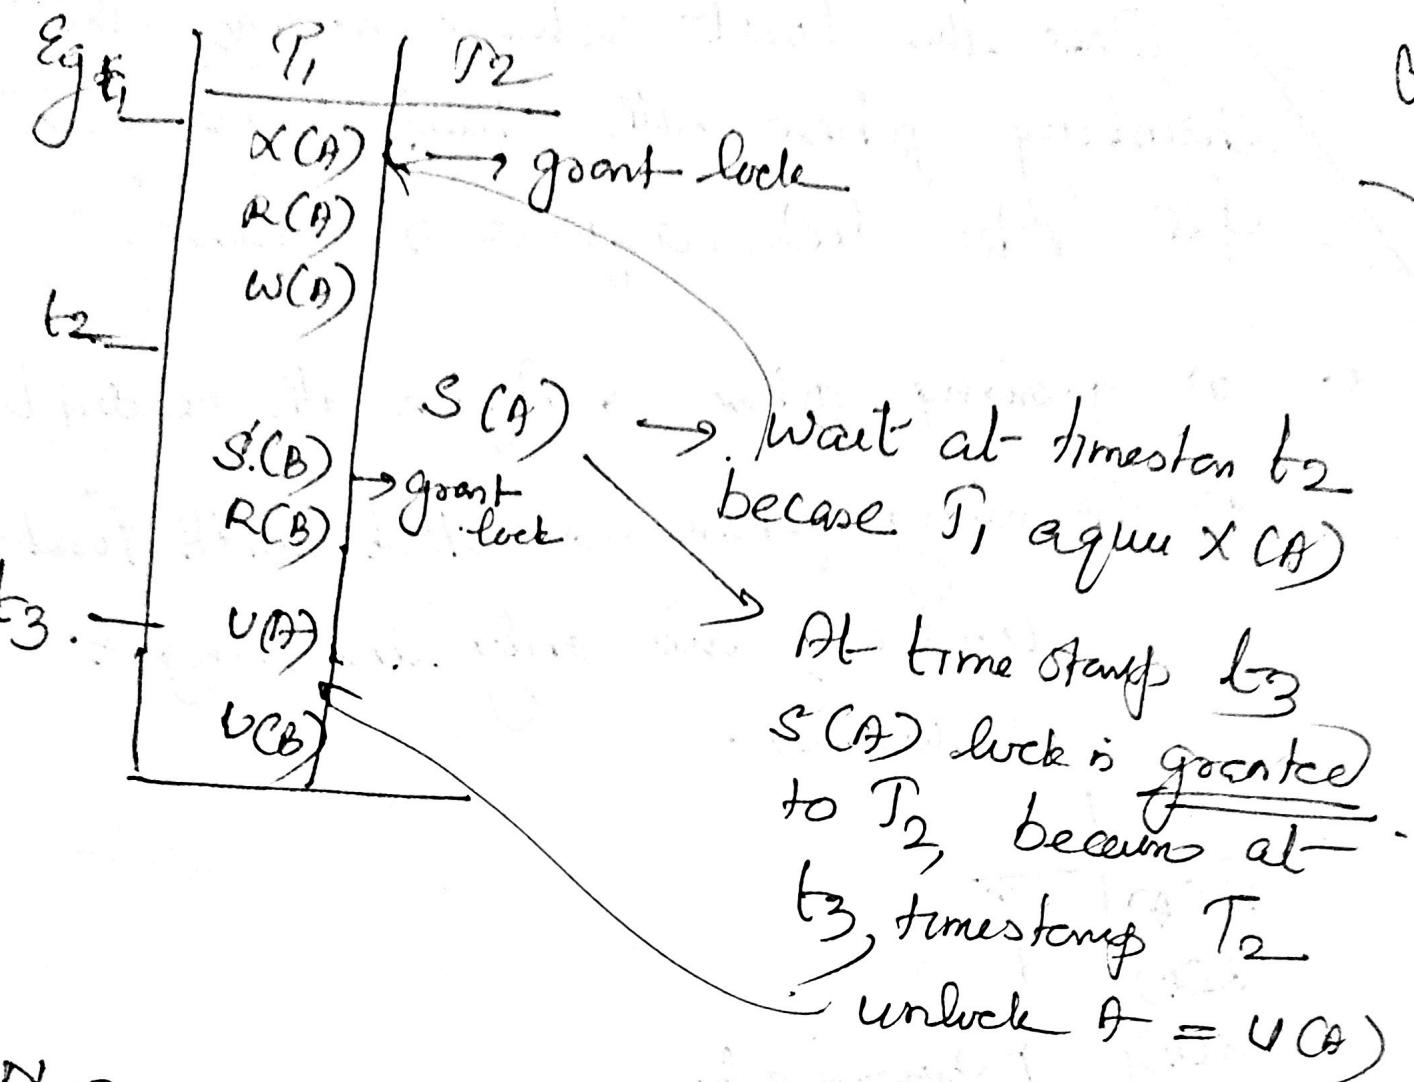
we acquire and release locks whenever needed.  
this continues till the end of transaction.

(ii) 2 PL, growing phase work grows as we continue acquiring locks. This phase ends when the first release occurs.

Once the first release occurs, the shrinking phase will start, and then after 10 lock acquire is allowed.

- a. in growing phase → Acquire all needed lock  
 ↳ shrinking phase → Start. with first unlock, and only unlocking is permitted.





Now

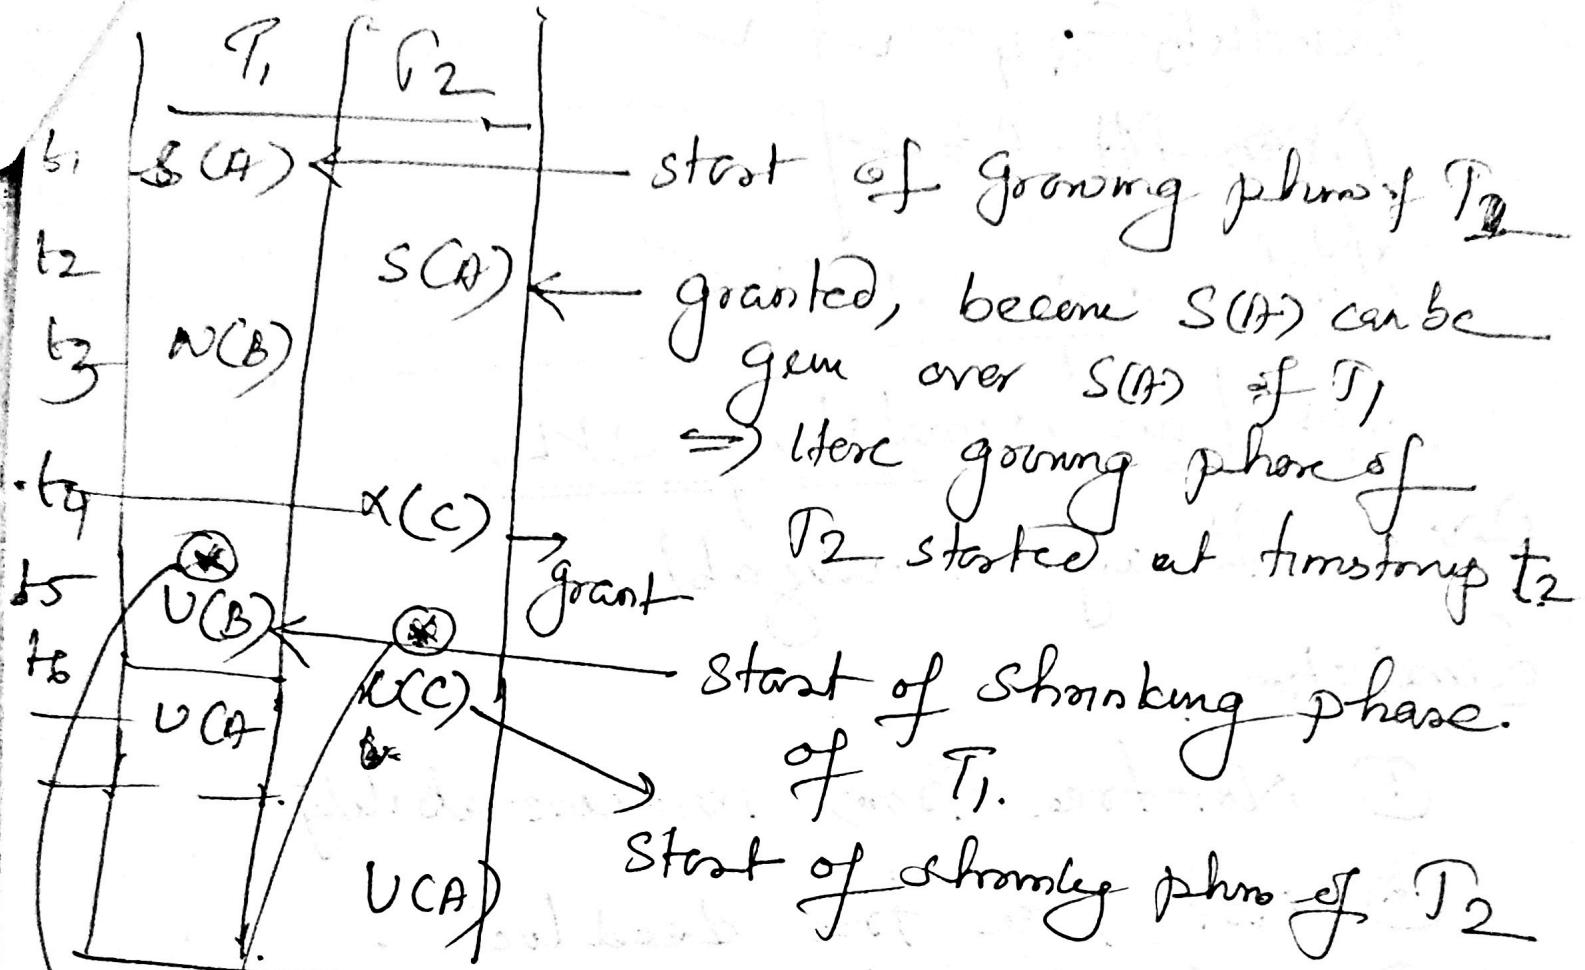
- ① it is achieving serializability. cohess was not ensured by SIX locking protocol. Here  $T_1$  execute like time stamp  $t_3$  and  $T_2$  starts after timestamp  $t_3$

$T_1 \rightarrow T_2 \Rightarrow$  No-loop.

Hence serializable.

Note:- Any transaction is running under 2PL protocol, it always serializable.

Eg-2



Here is the lock point of  $T_1$ , i.e. the time when shrinking phase starts.

Here is the lock point of  $T_2$ , i.e. the time when shrinking phase of  $T_2$  starts.

~~AT~~ S

Devdutt 24 ✓ ✓

Ajay P.D. 63 ✓

Nishant Rajeev ✓ ✓

Advantages and Limitations of 2PL

Adv : - Always serializable.

Limitations

- ① Not free from unrecoverability.
- ② Not free from dead lock.
- ③ Not free from starvation.
- ④ Not free from Cascading Roll back  
OR

①

$T_1$	$T_2$
X(A)	
R(A)	
W(A)	
V(A)	
	S(A)
	G(A)
	R(C)
	V(C)
Rollback fail	correct

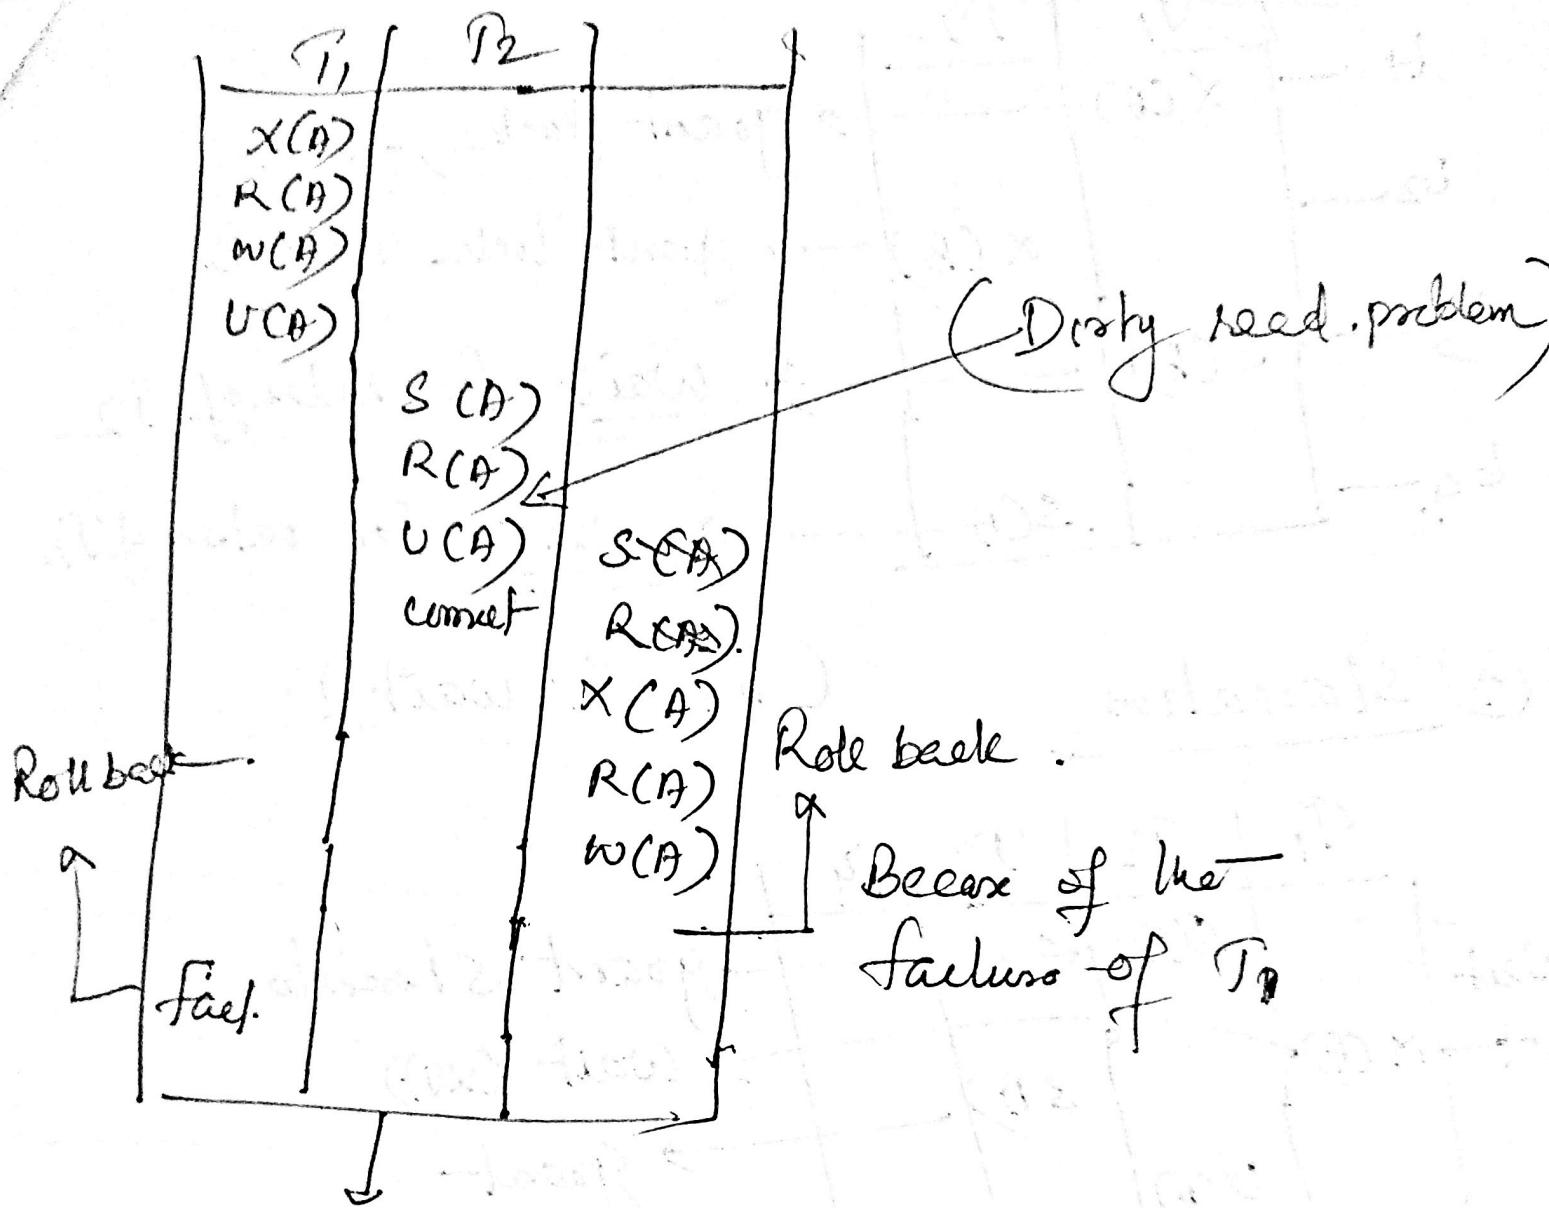
Similar to 2PC example.

$T_2$  committed first,

$T_1$  fails,  $T_1$  roll back

But  $T_2$  already committed so not able to rollback.

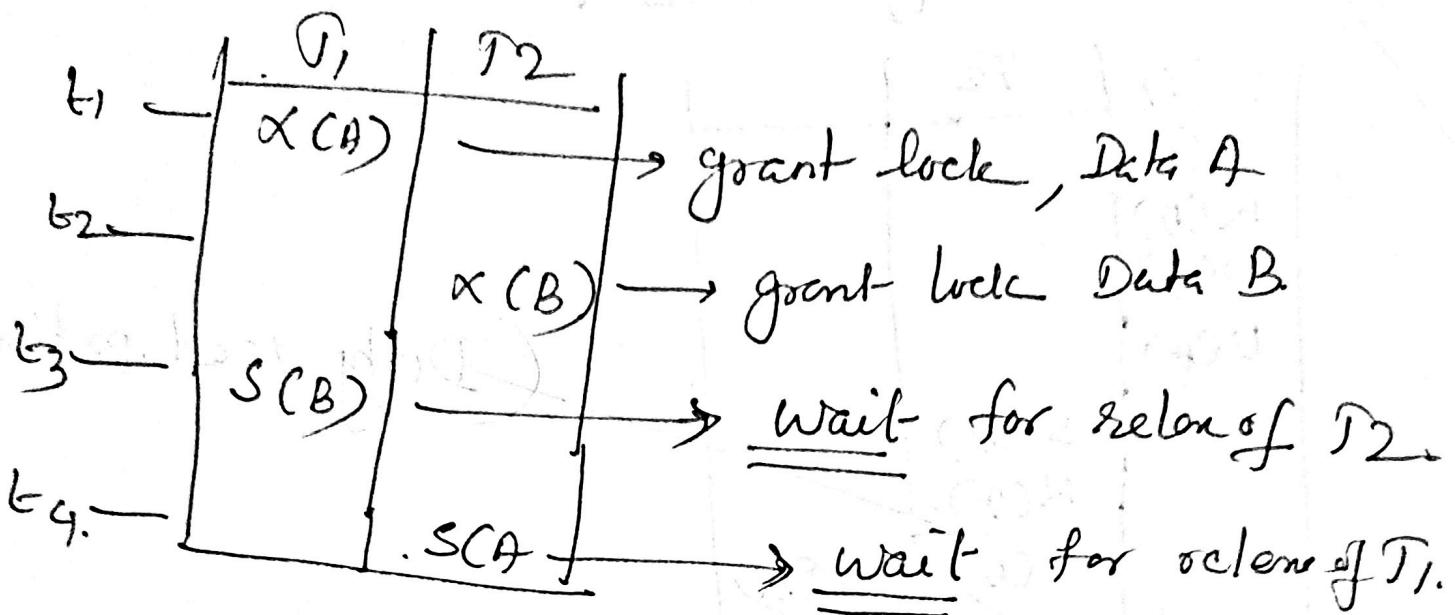
## ④ Cascading roll back



$T_2$  - Not possible to roll back (Dirty R  
Irrecoverable)

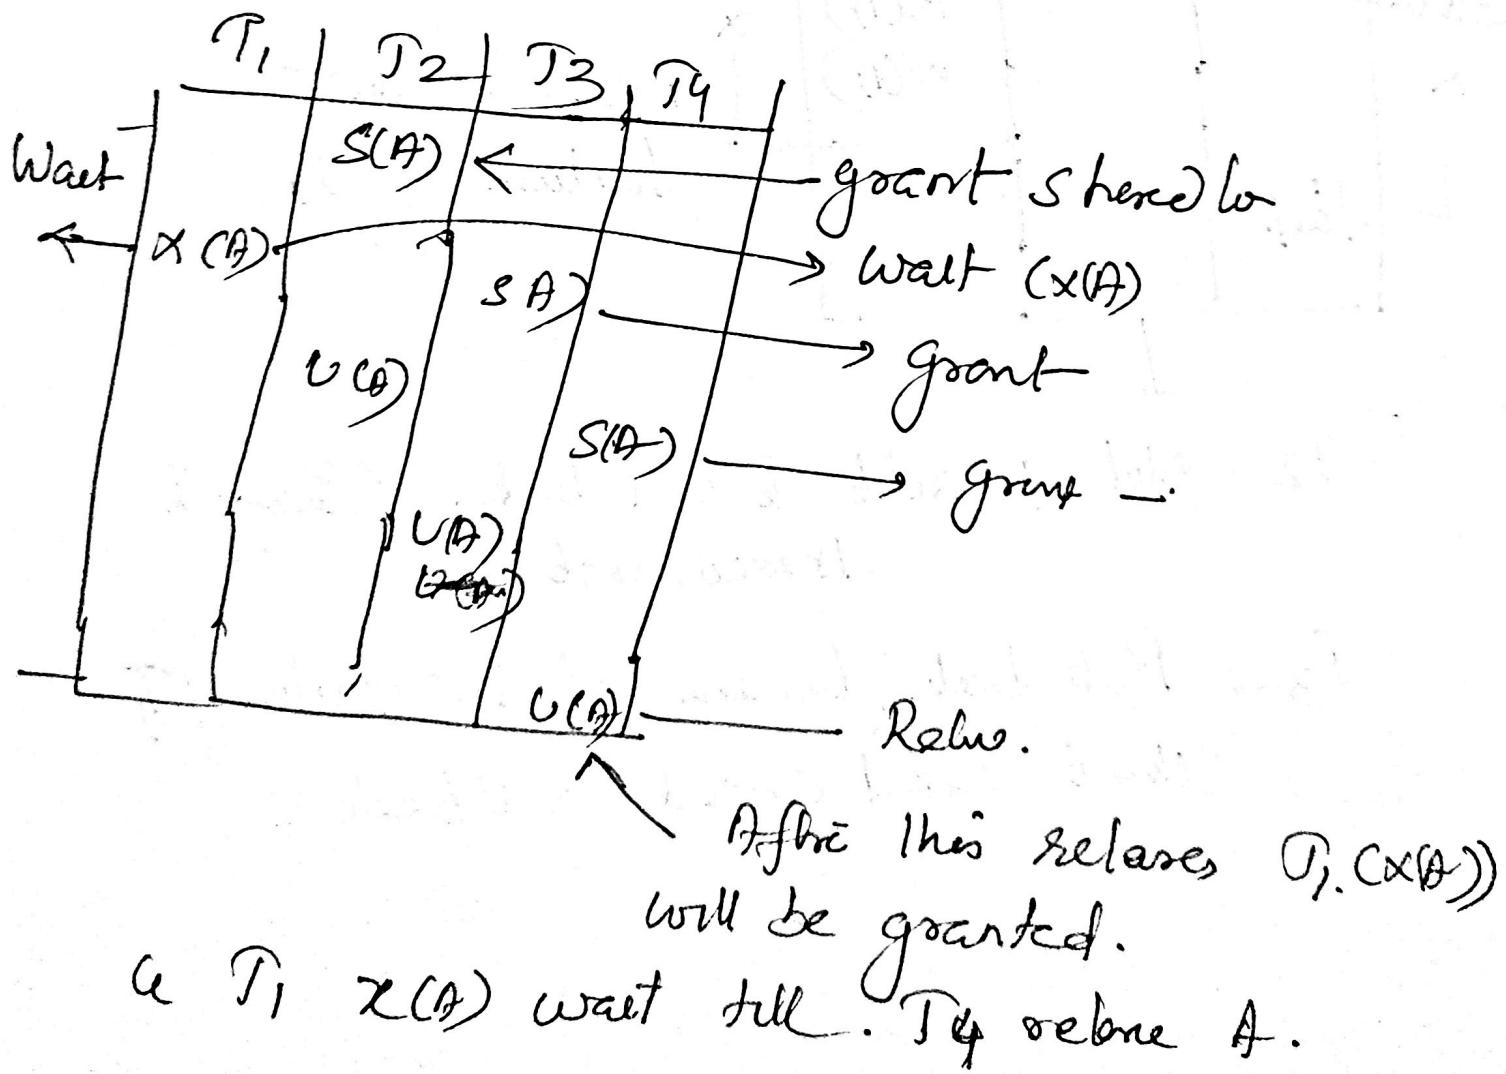
$T_3$  - Roll back because of the failure of  $T_2$   
This is called cascade rollback.

### ③ Dead lock.



### ④ Starvation

(infinite wait).



## Strict 2PL and Rigorous 2PL

Base 2PL has 4 limitations.

- ① Irrecoverability
- ② Cascading rollback
- ③ Deadlock
- ④ Starvation

Strict 2PL fix ① and ②

To achieve this, Certain restrictions are added to Base 2PL. That is

for Strict 2PL  $\Rightarrow$  it should satisfy Base 2PL and all exclusive locks should hold until commit / Abort

$T_1$	$T_2$	$T_3$	
$(X)A$			$T_1$ fails, and $T_2$ and $T_3$ also roll back. This is cascading rollback.
$R(A)$			
$W(A)$			
	$S(A)$		wait
	$R(A)$		$\rightarrow$ At this point $T_2 S(A)$ is granted. This also causes $T_3 (S)$ granted.
<u>Rollback</u>	$S(A)$	$R(A)$	
<u><math>U(A)</math></u>			
<u>fail</u>			

7) In between & after commit  
then  $T_2$ .

Here  $T_2$  and  $T_3$  are not reading from  $DB$  because,  $T_1$  has not committed yet.

$\Rightarrow$  If  $T_1$  unlock after Commit,  $T_2$  and  $T_3$  will read from  $DB$ , and can be. Roll-back under roll-back is not needed because of  $T_1$  failure.

	$T_1$	$T_2$	$T_3$	
$t_3$	$X(A)$			
	$R(A)$			
	$W(B)$			
$t_4$		$X(A)$		
		$R(A)$		
		$W(A)$		
$t_5$	$C$		$U(A)$	

$T_1$  commit at timestamp  $t_4$ , followed by  $U(C)$   
so at timestamp  $t_5$

$T_2$ .  $X(A)$  is granted

Now  $T_2$  is reading from  $DB$  not, that's not the value written by  $T_1$  in log.

Hence no need to rollback for  $T_2$ . This eliminating Cascaded rollback problem.

## remove irrecoverability

Since  $T_1$  unlock will happen at commit or abort. So  $T_2$  lock granted after hearing  $T_1$  commit.

(e) The first one to write commit first. This remove irrecoverability.

So Strict 2PL or,

- ① Always cascades.
- ② Recoverable.

Strict 2PL = Basic 2PL +  
hold exclusive lock until  
it commit / Abort

## Rigorous 2PL

[definition - Dead lock, Strict]

- ① Satisfy Strict 2PL + Hold shared lock until transaction commit / Abort

[That's the restoration is imposed on shared lock also]

## Rigorous 2PL

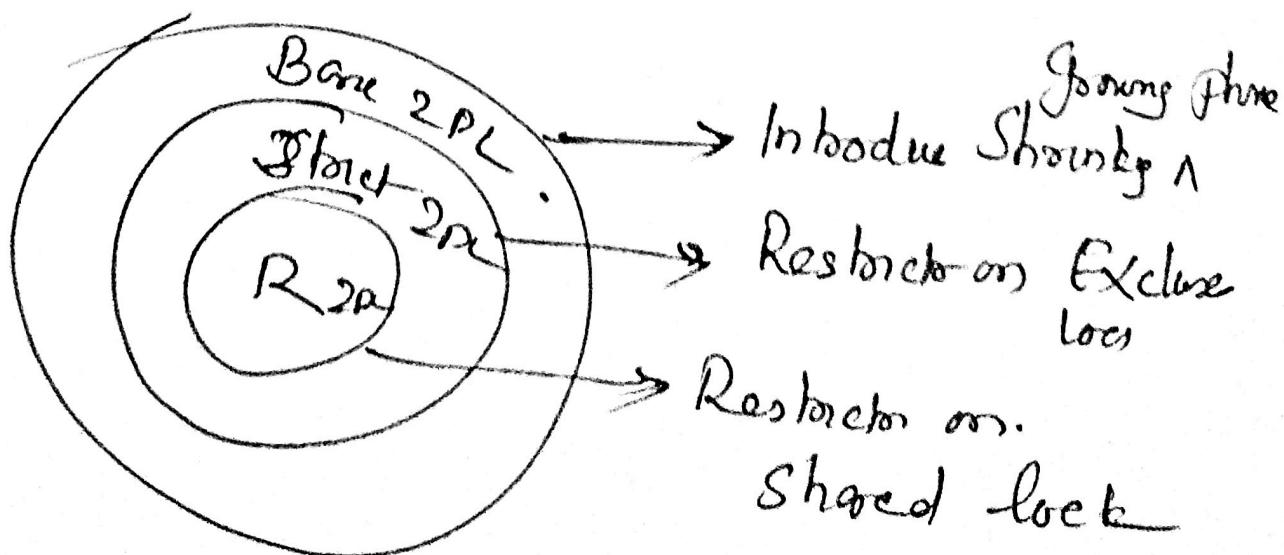
- ① Must support strict 2PL
- ② Restrictions applied to shared locks  
as shared lock is also released after commit

This does not ensure fix

- ① Deadlock
- ② Starvation

This fixes:

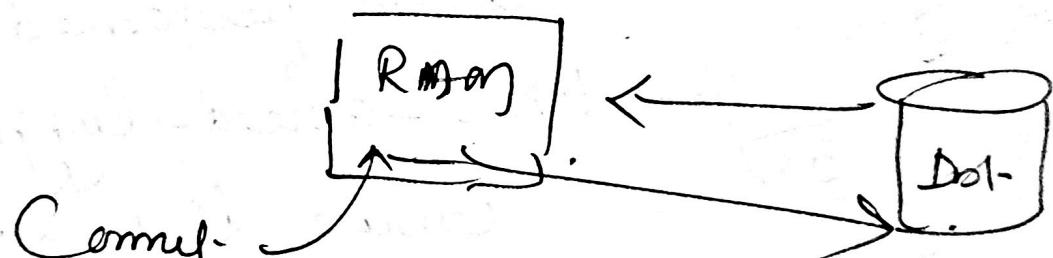
- ① Recoverability
- ② Cascade rollback



## Failure classifications

- ① → Logical Errors / Transaction failure.
  - Transaction can not proceed because of bad input, or overflow etc.
  - May be re-executed
- ② → System crash.
  - H/W malfunction
  - loss of RAM contents

### Transactions Power



Data loss does not occur

- ③ Disk failure.
  - Loss of Data
  - fix - Multiple copies of data in Backup disk.

To ensure consistency and atomicity, there will be an algorithm to recover from failure.

# How to Recover from failure.

①

Using an algorithm, run on failure case to perform recovery.

This Algo will have 2 parts -

- ① Action at normal transaction
- ② Action at failure (After failure)

① Those actions include

- ① - logging - Maintains a log that records the sequence of actions performed in transaction.
- ② - Write-ahead logging - (WAL)

Ensure that all log records associated with a transaction are written into stable storage (non volatile) before the pages' updated into Disk.

This ensure log records are durable on failure.

- ③ Checkpointing → Create checkpoints in log. Which will act as reference for recovery.

During checkpoint: DBMS writes all data pages to disk, and make their references on checkpoint in log.

## ② Action at failure.

- ① - Analysis - Analyse the log
- ② - Redo - Redo transactions that are committed but not updated to Disk
- ③ - Undo - Transactions which are not committed are rollbacked.  
(Partially completed transactions are rollbacked to ensure consistency).
- Restart - After redo and undo done session restarts the system & resume.

So combining ① Action at transaction and ② Action after failure, a recovery algorithm ensures database to restore to a consistent state.

## Log Based Recovery

Eg - I tried to withdraw 5000, and transaction failed at the end step.

My net balance should not be affected.  
The recovery is based on logs.

The type of log records are .

$\langle T_i^o \text{ Start} \rangle$   $T_i^o$  started .

$\langle T_i^o, x_j^o, V_1, V_2 \rangle$   $T_i^o$  performed work on data  $x_j^o$ . The initial val. value of

$\langle T_i^o \text{ Commit} \rangle$   $x_j^o$  before transaction is  $V_1$ .  
Value after transaction  $V_2$ .

$\langle T_i^o \text{ Abort} \rangle$   
( The purpose is to  
find the old value for  
roll back .

The older value  $\rightarrow V_1$

The new value  $V_2$

If transaction fails it will be rollbacked to  $V_1$

## Database Modifications

- ① Log is created prior to DB modification
- ② Log Record enables below on Recovery.
  - ① Redo  $\Rightarrow$  Committed but not updated to DB • Redo that transaction
  - ② Undo  $\Rightarrow$  Not committed.

(Partially committed)  
This enforces consistency (Completed).

Undo Transaction: performed when Tr is aborted / failed, and those are not committed  
Redo Transaction, those are committed but not updated into Disk

For both Redo, Undo, log is used to identify transaction, status, and ~~or~~ values.  
(Older values to rollback)

- ③ Once Redo / undo is performed, the data is updated to Disk.

## Deferred Modification

Date is written into disk after committing transaction.

One transaction has generally one commit. So deferred modification updates at the end of transaction. It will not write any intermediate result values.

## Immediate Modification

- DB is updated in between transactions.
- Not waiting for commit or intentions before committing.

## Cheek point - Checkpointing

Another type of entry is cheek point. It is just like other records; it is an entry in log similar to other type of records we found earlier.

Cheek point is a log indicate at that timestamp point, the system writes out all DBMS buffers that have been modified to disk.

So essentially at check point all data is written into disk. It comes Checkpoint comes periodically in "immediate modification".

It is really a checkpoint record, rather than a reference in log file.

The checkpoint record holds Transaction IDs of active transaction those have done database write at that point. So using checkpoint, system can identify the transaction during recovery.

<checkpoint-  $T_1 T_2 T_2 \dots T_3$  >  
Transactions IDs.

Sequence of Actions in Checkpointing.

Steps performing during checkpointing.

Checkpointing → is a process of taking checkpoints by DBMS at specific intervals or when specific conditions are met.

Checkpointing involve

① Determine checkpoint timing.

(When to initiate checkpointing)

May be based on elapsed time,  
no of transactions processed etc.

② Trigger checkpoint - Once conditions are met, DBMS initiate checkpoint process.

③ Execute checkpoint Action -

→ flush data to disk

→ Write a checkpoint record in log

→ Update checkpoint information.

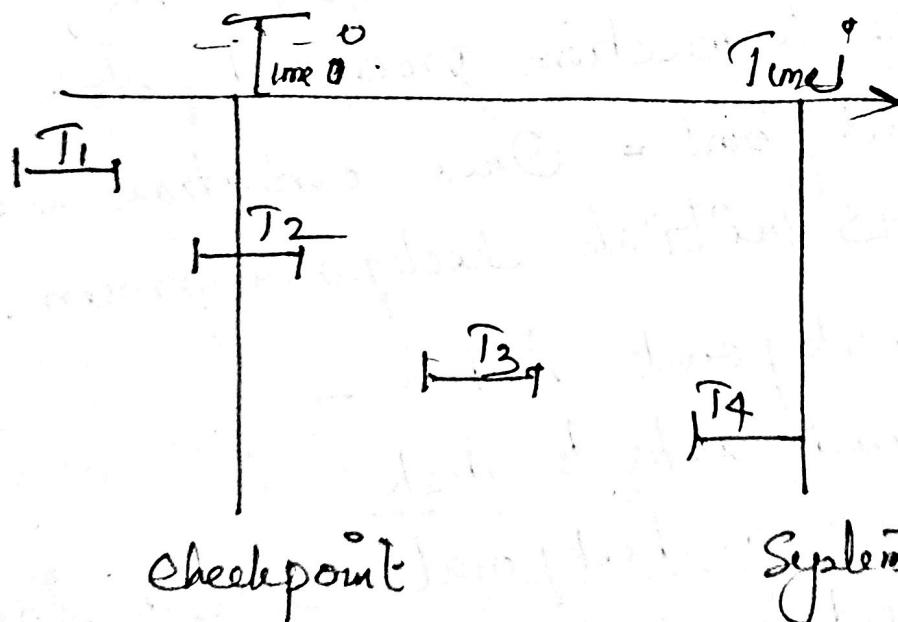
(Transaction id etc).

④ Resume normal operation.

[Checkpointing can be Time based, Transaction Based,  
etc]

- 4
- ④ Transaction are not allowed during checkpointing to perform any db write or writing to log file during checkpointing
  - ⑤ After system crash occurs, the DBMS system scans log file to locate the last checkpoint entry (Checkpoint)

### Recovery Procedure



$T_4$  - Started after checkpoint but not committed before failure.

$T_1$  - Already committed,  
(Before checkpoint)

System Failure .

$T_1, T_2, T_3, T_4$  are  
Transactions .

$T_2$  - Committed but not updated

$T_3$  - Committed but not updated

$T_4$  - Not committed

## Recovery.

Redo  $\rightarrow T_2 \& T_3$ , Because committed but not updated to disk..

Update to disk is done at <sup>next</sup> checkpoints, which is expected to be taken by DBMS ~~next~~.

$\rightarrow$  So redo  $T_2$  and  $T_3$

Undo  $\rightarrow T_4$ , which is not committed.

i.e. Partially completed transaction

so undo  $T_4$ .

$T_1 \rightarrow$  can be ignored, as it is committed before checkpoint and updated to disk at the given checkpoint.  
 $\rightarrow$  So it can be ignored.

Algo: During recovery need to consider only the most recent transactions; that started before checkpoint

① Scan backwards from failure to find the last checkpoint. (Most recent checkpoint)  
 $\langle$ checkpoint,  $t$  $\rangle$

② Continue scanning backward to get  
 $\langle T_1 \text{ start} \rangle$  in log.

- ③ Need only consider transaction started before check point
- ④ for all transactions with NO  $\langle$  Commit  $\rangle$   
execute undo( $T_i$ )
- ⑤ All transactions with a  $\langle$   $T_i$ , Commit  $\rangle$   
execute redo( $T_i$ )

These are the cases for immediate modification where check disk update is done periodically.

Next Recovery based on deferred update

- Deferred Modification  $\rightarrow$  Data is updated to Disk after the final commit.
- Deferred update postpone any actual update to the database on disk until the transaction commits.
- After commit, the value in log is written to DB.

Here UNDO is NOT required, because, the transaction failed before it commits and wrote data onto DB in Disk. Since no DB update on Disk is performed, nothing to rollback or undo.

Normal Case  $\rightarrow$  When checkpoint is reached.

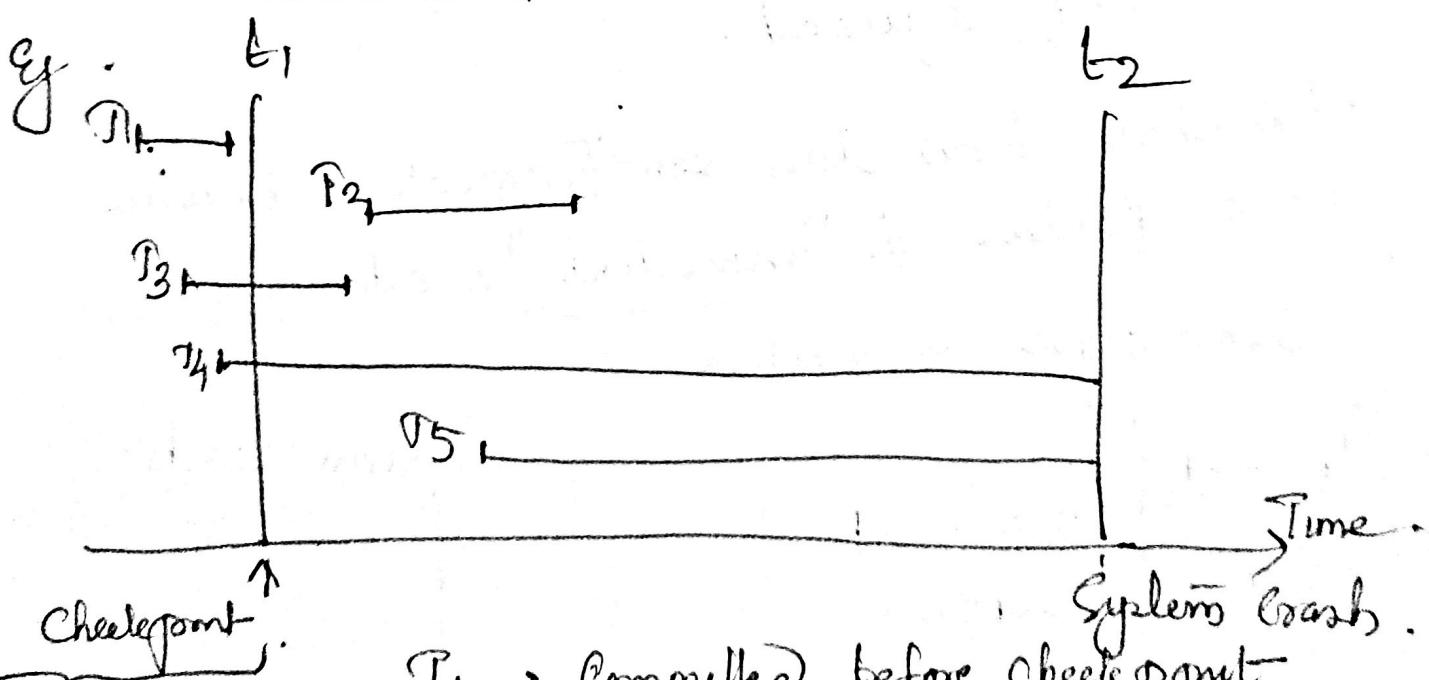
DBMS writes all data into DB in Disk.

Abort Case  $\rightarrow$  Crashes before checkpoint.

Ideally all committed transactions write data into DB in Disk at checkpoint.

So if crashes before checkpoint, REDO all committed transactions; to be precise all WRITE operations of all committed transaction.

There can be uncommitted transactions. They are effectively cancelled and must be resubmitted or reexecuted.



Some Transactions  
Committed are  
Taken checkpoint  
to write DB

$T_1 \rightarrow$  Committed before checkpoint

$T_2, T_3 \rightarrow$  Not committed but not written to DB in Disk

$T_4, T_5 \rightarrow$  Active Transaction, that is not committed.

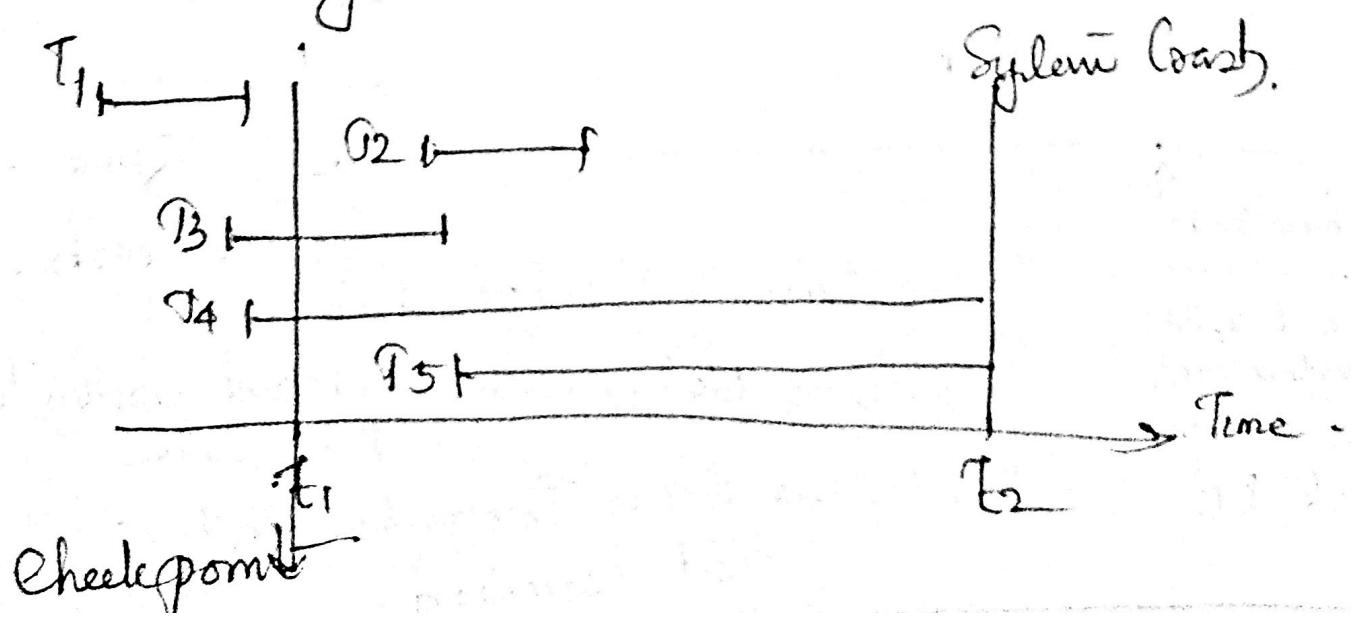
## Actions

$T_1 \rightarrow$  Already Committed and written to Disk at checkpoint.

$T_2, T_3 \rightarrow$  Committed after checkpoint -  
not written to DB is disk Before that  
the system crashed  
So need to RERDO.

$T_4, T_5 \rightarrow$  Not committed, active transactions.  
Since not committed, database is not updated.. Hence ignore.  
So it may be resubmitted or reexecuted if required.

Now check how these transactions behave  
on a failure in "immediate" modification.  
Same figure as earlier



$T_1$ : Already committed before checkpoint, and at checkpoint update DB in Disk.

$T_2$ : Committed after checkpoint, but not written data into DB in Disk.

$T_3$ : Started before checkpoint, and at checkpoint, some write operations may have written to DB in Disk.

$T_4$ : Similar to  $T_3$ , but  $T_4$  is not committed.

$T_5$ : Database Not committed, (Because nothing is written to disk because started after checkpoint).

$T_1$ : Ignore.

$T_2$ : Redo since it is committed.

$T_3$ : Redo " "

$T_4$  &  $T_5$ ! - Not committed, Before commit, it failed, undo (Rollback)